

# Music Classification

July 8, 2019

```
In [21]: import warnings
warnings.filterwarnings("ignore", category = FutureWarning)

import os
import keras
import h5py
import librosa
import itertools
import numpy as np
import matplotlib.pyplot as plt

# Preprocessing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler

from keras.utils import to_categorical
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix

from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Activation
from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.layers import Dropout
from keras.layers import Flatten
from keras.layers import BatchNormalization
```

## 1 Loading the Data

### 1.0.1 Defining general functions

```
In [34]: # Function to split a single song into overlapping windows of the same
import more_itertools as it_plus

# def split_window(arr, genre, chunk_size=0.1, step_size=1):
```

```

#         # Get the window size
#         window_size = int(X.shape[0]*chunk_size)
#         final_arr_x = list(it_plus.windowed(X, n=window_size, step=step_size))
#         final_arr_y = [genre for x in range(len(final_arr_x))]
#         print(final_arr_x)
#         return np.array(final_arr_x), np.array(final_arr_y)

def split_window(X, y, window = 0.1, overlap = 0.5):
    # Empty lists to hold our results
    temp_X = []
    temp_y = []

    # Get the input song array size
    xshape = X.shape[0]
    chunk = int(xshape*window)
    offset = int(chunk*(1.-overlap))

    # Split the song and create new ones on windows
    spsong = [X[i:i+chunk] for i in range(0, xshape - chunk + offset, offset)]
    for s in spsong:
        temp_X.append(s)
        temp_y.append(y)

    return np.array(temp_X), np.array(temp_y)

```

In [33]: #Function to convert songs into their respective Melspectograms  
[#https://github.com/Hguimaraes/gtzan.keras/blob/master/src/gtzan/struct.py](https://github.com/Hguimaraes/gtzan.keras/blob/master/src/gtzan/struct.py)

```

def to_melspec(song, n_fft=1024, hop_len=512):
    #Transform into spectrograms
    melspec = lambda x: librosa.feature.melspectrogram(x,n_fft = n_fft,hop_length = hop_len)
    transformed_songs = map(melspec, song)
    return np.array(list(transformed_songs))

```

In [36]: [#https://github.com/Hguimaraes/gtzan.keras/blob/master/src/gtzan/struct.py](https://github.com/Hguimaraes/gtzan.keras/blob/master/src/gtzan/struct.py)

```

def read_data(src_dir, genres, song_samples, spec_format, chunk_size=0.1, step_size=0.1):
    # Empty array of dicts with the processed features from all files
    arr_specs = []
    arr_genres = []

    # Read files from the folders
    for x, _ in genres.items():
        folder = src_dir + x

        for root, subdirs, files in os.walk(folder):
            for file in files:
                # Read the audio file

```

```

file_name = folder + "/" + file
signal, sr = librosa.load(file_name)
signal = signal[:song_samples]

# Debug process
if debug:
    print("Reading file: {}".format(file_name))

# Convert to dataset of spectrograms/melspectrograms
signals, y = split_window(signal, genres[x])

# Convert to "spec" representation
specs = spec_format(signals)

# Save files
arr_genres.extend(y)
arr_specs.extend(specs)

return np.array(arr_specs), np.array(arr_genres)

```

```
In [37]: print(os.getcwd()+"/genres/")
```

E:\Work\Case Study\Music Classification\genres/

```

In [38]: if not (os.path.isfile("x_gtzan_npy.npy") or os.path.isfile("y_gtzan_npy.npy")):
# Parameters
gtzan_dir = os.getcwd()+"/genres/"
song_samples = 660000
genres = {'metal': 0, 'disco': 1, 'classical': 2, 'hiphop': 3, 'jazz': 4,
          'country': 5, 'pop': 6, 'blues': 7, 'reggae': 8, 'rock': 9}

# Read the data
X, y = read_data(gtzan_dir, genres, song_samples, to_melspec, debug=False)
np.save('x_gtzan_npy.npy', X)
np.save('y_gtzan_npy.npy', y)

```

```

In [41]: X = np.load('x_gtzan_npy.npy')
y = np.load('y_gtzan_npy.npy')

```

## 1.1 Building DL models

```

In [46]: import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error

```

```
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [42]: # One hot encoding of the labels
        y = to_categorical(y)
```

```
In [44]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=)
```

### 1.1.1 CNN

```
In [134]: from keras.models import Sequential
        from keras.layers import LSTM, BatchNormalization, TimeDistributed, Input
        from keras.layers.core import Dense, Dropout
        import keras
```

```
In [151]: # Model Definition
        timesteps = len(X_train[0])
        input_dim = len(X_train[0][0])
        print(timesteps)
        print(input_dim)
        print(X_train.shape)
        num_genres = 10
```

128

129

(13300, 128, 129, 1)

```
In [179]: # inp_shape = Input(shape=X_train[0].shape)
        # print(inp_shape)
        # #Model 1
        # model_1 = Conv2D(64, (4,4), padding='same', activation='relu')(inp_shape)
        # model_1 = MaxPooling2D((2,2), strides=(1,1), padding='same')(model_1)
        # model_1 = Dropout(0.5)(model_1)
        # model_1 = Conv2D(128, (4,4), padding='same', activation='relu')(model_1)
        # model_1 = MaxPooling2D((2,2), strides=(1,1), padding='same')(model_1)
        # # model_1 = Flatten()(model_1)
        # # model_1 = Dropout(0.8)(model_1)

        # #Model 2
        # model_2 = LSTM(64, kernel_initializer=keras.initializers.glorot_normal(seed=None),
        # model_2 = Dropout(0.8)(model_2)
        # model_2 = LSTM(128, kernel_initializer=keras.initializers.glorot_normal(seed=None),
        # model_2 = Dropout(0.6)(model_2)
        # merged = keras.layers.concatenate([model_1, model_2], axis=1)
```

```

# out = Dense(64, activation='relu')(merged)
# out = Dropout(0.7)(merged)
# out = Dense(num_genres, activation='softmax')(out)
# model = Model(input_shape, out)
# model.summary()
# model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

```

In [192]: # Model Definition

```

input_shape = X_train[0].shape
num_genres = 10

model = Sequential()
# Conv Block 1
model.add(Conv2D(16, kernel_size=(3, 3), strides=(1, 1),
                 activation='relu', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Dropout(0.25))

# Conv Block 2
model.add(Conv2D(32, (3, 3), strides=(1, 1), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Dropout(0.25))

# Conv Block 3
model.add(Conv2D(64, (3, 3), strides=(1, 1), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Dropout(0.25))

# Conv Block 4
model.add(Conv2D(128, (3, 3), strides=(1, 1), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Dropout(0.4))

# Conv Block 5
model.add(Conv2D(64, (3, 3), strides=(1, 1), activation='relu'))
model.add(MaxPooling2D(pool_size=(4, 4), strides=(4, 4)))
model.add(Dropout(0.3))

# MLP
model.add(Flatten())
model.add(Dense(num_genres, activation='softmax'))

model.summary()

```

Layer (type)	Output Shape	Param #
conv2d_127 (Conv2D)	(None, 126, 127, 16)	160

```

-----
max_pooling2d_100 (MaxPoolin (None, 63, 63, 16)          0
-----
dropout_104 (Dropout)          (None, 63, 63, 16)          0
-----
conv2d_128 (Conv2D)            (None, 61, 61, 32)          4640
-----
max_pooling2d_101 (MaxPoolin (None, 30, 30, 32)          0
-----
dropout_105 (Dropout)          (None, 30, 30, 32)          0
-----
conv2d_129 (Conv2D)            (None, 28, 28, 64)          18496
-----
max_pooling2d_102 (MaxPoolin (None, 14, 14, 64)          0
-----
dropout_106 (Dropout)          (None, 14, 14, 64)          0
-----
conv2d_130 (Conv2D)            (None, 12, 12, 128)          73856
-----
max_pooling2d_103 (MaxPoolin (None, 6, 6, 128)           0
-----
dropout_107 (Dropout)          (None, 6, 6, 128)           0
-----
conv2d_131 (Conv2D)            (None, 4, 4, 64)            73792
-----
max_pooling2d_104 (MaxPoolin (None, 1, 1, 64)            0
-----
dropout_108 (Dropout)          (None, 1, 1, 64)            0
-----
flatten_48 (Flatten)           (None, 64)                   0
-----
dense_48 (Dense)               (None, 10)                   650
=====
Total params: 171,594
Trainable params: 171,594
Non-trainable params: 0
-----

```

```
In [193]: X_train[0].shape
```

```
Out[193]: (128, 129, 1)
```

```
In [194]: model.compile(loss=keras.losses.categorical_crossentropy,
                        optimizer=keras.optimizers.Adam(),
                        metrics=['accuracy'])
```

```
In [195]: hist1 = model.fit(X_train, y_train,
                           batch_size=128,
```

```
epochs=40,  
verbose=1,  
validation_data=(X_test, y_test))
```

Train on 13300 samples, validate on 5700 samples

```
Epoch 1/40  
13300/13300 [=====] - 99s 7ms/step - loss: 2.3367 - acc: 0.1909 - val.  
Epoch 2/40  
13300/13300 [=====] - 98s 7ms/step - loss: 2.0150 - acc: 0.2623 - val.  
Epoch 3/40  
13300/13300 [=====] - 101s 8ms/step - loss: 1.8090 - acc: 0.3523 - val.  
Epoch 4/40  
13300/13300 [=====] - 101s 8ms/step - loss: 1.6439 - acc: 0.4183 - val.  
Epoch 5/40  
13300/13300 [=====] - 101s 8ms/step - loss: 1.5231 - acc: 0.4613 - val.  
Epoch 6/40  
13300/13300 [=====] - 102s 8ms/step - loss: 1.4365 - acc: 0.4942 - val.  
Epoch 7/40  
13300/13300 [=====] - 102s 8ms/step - loss: 1.3813 - acc: 0.5115 - val.  
Epoch 8/40  
13300/13300 [=====] - 102s 8ms/step - loss: 1.3043 - acc: 0.5430 - val.  
Epoch 9/40  
13300/13300 [=====] - 102s 8ms/step - loss: 1.2677 - acc: 0.5605 - val.  
Epoch 10/40  
13300/13300 [=====] - 102s 8ms/step - loss: 1.2190 - acc: 0.5804 - val.  
Epoch 11/40  
13300/13300 [=====] - 100s 7ms/step - loss: 1.1745 - acc: 0.5956 - val.  
Epoch 12/40  
13300/13300 [=====] - 93s 7ms/step - loss: 1.1474 - acc: 0.6098 - val.  
Epoch 13/40  
13300/13300 [=====] - 94s 7ms/step - loss: 1.1025 - acc: 0.6258 - val.  
Epoch 14/40  
13300/13300 [=====] - 99s 7ms/step - loss: 1.0709 - acc: 0.6330 - val.  
Epoch 15/40  
13300/13300 [=====] - 95s 7ms/step - loss: 1.0466 - acc: 0.6473 - val.  
Epoch 16/40  
13300/13300 [=====] - 94s 7ms/step - loss: 1.0117 - acc: 0.6562 - val.  
Epoch 17/40  
13300/13300 [=====] - 96s 7ms/step - loss: 0.9957 - acc: 0.6601 - val.  
Epoch 18/40  
13300/13300 [=====] - 95s 7ms/step - loss: 0.9696 - acc: 0.6708 - val.  
Epoch 19/40  
13300/13300 [=====] - 97s 7ms/step - loss: 0.9599 - acc: 0.6752 - val.  
Epoch 20/40  
13300/13300 [=====] - 95s 7ms/step - loss: 0.9178 - acc: 0.6947 - val.  
Epoch 21/40  
13300/13300 [=====] - 95s 7ms/step - loss: 0.9079 - acc: 0.6964 - val.  
Epoch 22/40
```

```

13300/13300 [=====] - 95s 7ms/step - loss: 0.8863 - acc: 0.7019 - val.
Epoch 23/40
13300/13300 [=====] - 95s 7ms/step - loss: 0.8728 - acc: 0.7053 - val.
Epoch 24/40
13300/13300 [=====] - 96s 7ms/step - loss: 0.8519 - acc: 0.7158 - val.
Epoch 25/40
13300/13300 [=====] - 96s 7ms/step - loss: 0.8343 - acc: 0.7222 - val.
Epoch 26/40
13300/13300 [=====] - 97s 7ms/step - loss: 0.7856 - acc: 0.7359 - val.
Epoch 27/40
13300/13300 [=====] - 98s 7ms/step - loss: 0.7815 - acc: 0.7388 - val.
Epoch 28/40
13300/13300 [=====] - 97s 7ms/step - loss: 0.7875 - acc: 0.7335 - val.
Epoch 29/40
13300/13300 [=====] - 96s 7ms/step - loss: 0.7692 - acc: 0.7441 - val.
Epoch 30/40
13300/13300 [=====] - 99s 7ms/step - loss: 0.7500 - acc: 0.7495 - val.
Epoch 31/40
13300/13300 [=====] - 97s 7ms/step - loss: 0.7386 - acc: 0.7533 - val.
Epoch 32/40
13300/13300 [=====] - 98s 7ms/step - loss: 0.7459 - acc: 0.7502 - val.
Epoch 33/40
13300/13300 [=====] - 98s 7ms/step - loss: 0.7254 - acc: 0.7595 - val.
Epoch 34/40
13300/13300 [=====] - 100s 8ms/step - loss: 0.7072 - acc: 0.7618 - val.
Epoch 35/40
13300/13300 [=====] - 101s 8ms/step - loss: 0.7026 - acc: 0.7668 - val.
Epoch 36/40
13300/13300 [=====] - 100s 8ms/step - loss: 0.6870 - acc: 0.7678 - val.
Epoch 37/40
13300/13300 [=====] - 101s 8ms/step - loss: 0.6798 - acc: 0.7747 - val.
Epoch 38/40
13300/13300 [=====] - 100s 8ms/step - loss: 0.6663 - acc: 0.7796 - val.
Epoch 39/40
13300/13300 [=====] - 97s 7ms/step - loss: 0.6677 - acc: 0.7795 - val.
Epoch 40/40
13300/13300 [=====] - 97s 7ms/step - loss: 0.6670 - acc: 0.7759 - val.

```

```

In [197]: score = model.evaluate(X_test, y_test, verbose=0)
          print('Test score:', score[0])

          print('Test accuracy:', score[1])
          fig,ax = plt.subplots(1,1)
          ax.set_xlabel('epoch') ; ax.set_ylabel('Sparse Categorical Crossentropy Loss')
          # list of epoch numbers
          x = list(range(1,40+1))

```




```

vy = hist1.history['val_loss']
ty = hist1.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.5945673463846508

Test accuracy: 0.8049122807017544



### 1.1.2 Iterating for another 10 epochs

```

In [198]: hist2 = model.fit(X_train, y_train,
                             batch_size=128,
                             epochs=10,
                             verbose=1,
                             validation_data=(X_test, y_test))

```

Train on 13300 samples, validate on 5700 samples

Epoch 1/10

13300/13300 [=====] - 96s 7ms/step - loss: 0.6364 - acc: 0.7912 - val.

Epoch 2/10

13300/13300 [=====] - 96s 7ms/step - loss: 0.6333 - acc: 0.7910 - val.

Epoch 3/10

13300/13300 [=====] - 99s 7ms/step - loss: 0.6286 - acc: 0.7886 - val.

Epoch 4/10

```

13300/13300 [=====] - 97s 7ms/step - loss: 0.6219 - acc: 0.7922 - val_
Epoch 5/10
13300/13300 [=====] - 96s 7ms/step - loss: 0.6211 - acc: 0.7890 - val_
Epoch 6/10
13300/13300 [=====] - 98s 7ms/step - loss: 0.6103 - acc: 0.7923 - val_
Epoch 7/10
13300/13300 [=====] - 98s 7ms/step - loss: 0.6137 - acc: 0.7938 - val_
Epoch 8/10
13300/13300 [=====] - 98s 7ms/step - loss: 0.6045 - acc: 0.8006 - val_
Epoch 9/10
13300/13300 [=====] - 96s 7ms/step - loss: 0.5873 - acc: 0.8030 - val_
Epoch 10/10
13300/13300 [=====] - 98s 7ms/step - loss: 0.5833 - acc: 0.8038 - val_

```

```

In [199]: score = model.evaluate(X_test, y_test, verbose=0)
          print('Test score:', score[0])

          print('Test accuracy:', score[1])
          fig,ax = plt.subplots(1,1)
          ax.set_xlabel('epoch') ; ax.set_ylabel('Sparse Categorical Crossentropy Loss')
          # list of epoch numbers
          x = list(range(1,10+1))


          vy = hist2.history['val_loss']
          ty = hist2.history['loss']
          plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.5268087197186654
Test accuracy: 0.8278947368421052

```



```
In [200]: model.save("music_genre_classn.h5")
```

```
In [ ]:
```